

Smartphone-Only Estimation of Elevator Vertical Displacement with Calibrated Confidence Intervals

EYAL YAKIR, TBD, TBD

Knowing which floor a smartphone user is on remains an open problem on devices that lack a pressure sensor, and the single most common floor-change event — the elevator ride — is exactly where the three workhorses of indoor localisation each break down. Pedestrian dead reckoning and step counting estimate horizontal trajectory; the cabin removes the steps and the accelerometer cannot be integrated vertically without unbounded drift. Wi-Fi fingerprinting demands a per-floor survey, degrades whenever the radio environment drifts, and is shielded by the cabin itself. Barometric altimetry resolves a floor in principle but assumes a pressure sensor that entry-level phones omit, and demands re-referencing against weather drift on phones that carry one.

We present a smartphone-only pipeline that estimates the signed vertical displacement Δh of every elevator ride a user takes and attaches a calibrated 90 % confidence interval to it, using the phone’s accelerometer alone. The pipeline rests on one observation. The mechanical limits of an elevator cabin force the gravity-removed acceleration trace a phone records during a ride into a small parametric family of basis functions. We decompose every candidate ride against this family in two stages — segmentation of hours-long inertial streams into discrete up / down ride intervals, and per-ride prediction of Δh with a per-ride uncertainty. Each interval is distribution-free with finite-sample coverage guarantees, adapts to ride duration, gravity-vector stability, and fit quality, and remains valid under both white sensor noise and the coloured (drift- and vibration-driven) noise that dominates real hand-held recordings.

We evaluate on [TODO: 150] experiments spanning [TODO: TBD] elevators in [TODO: TBD] buildings (residential, office, hotel), with a strict held-out blind test. Median CI half-width on the held-out test stays below [TODO: ± 0.9 m] on short rides at 90 % coverage, an improvement of [TODO: TBD %] over the dataset’s prior best. We open-source the pipeline, dataset, and evaluation harness.

CCS Concepts: • **Human-centered computing** → **Ubiquitous and mobile computing**; • **Computing methodologies** → *Sensor data processing*.

Additional Key Words and Phrases: smartphone sensing, inertial measurement, elevator, vertical localization, conformal prediction, uncertainty quantification, indoor positioning, matched filter

ACM Reference Format:

Eyal Yakir. 2026. Smartphone-Only Estimation of Elevator Vertical Displacement with Calibrated Confidence Intervals. *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol.* 0, 0, Article 0 (2026), 31 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Indoor location services place a smartphone within metres of the right room on the horizontal plane, using Wi-Fi fingerprinting, BLE beacons, and visual pedestrian dead reckoning [9]. The vertical axis lags behind. Knowing the user’s floor matters for emergency dispatch, which routes responders to a floor rather than to a building entrance. It matters for venue-aware advertising and indoor delivery in malls and office towers, and for augmented-reality experiences anchored to a specific floor of a museum or conference centre.

Author’s Contact Information: Eyal Yakir, eyalyakir159@gmail.com, TBD, TBD, TBD.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

Three structural limitations make floor estimation harder than its horizontal sibling. (i) Step counting and pedestrian dead reckoning estimate horizontal trajectory and report nothing about height [5]. (ii) Wi-Fi fingerprinting separates floors only where the survey covered every floor and the radio environment has not drifted since [2]. (iii) Barometric altimetry resolves a floor in principle, but the pressure sensor is absent on entry-level phones and demands re-referencing against weather drift on phones that carry one [11]. The single most common floor-change event — the elevator ride — sits in the regime where every horizontal-localisation signal degrades at once.

Elevators are not arbitrary motion. We found that the mechanical limits of the cabin force the gravity-removed acceleration trace a phone records during a ride into a small parametric family of expected functions. Section 2 characterises this family precisely; what matters in the introduction is that the family is small enough to enumerate and structured enough to fit a candidate ride against. Every claim in this paper rests on it.

The family turns a hard goal into a tractable one. Given an inertial recording of a single elevator ride, we recover the ride’s signed vertical displacement Δh together with a calibrated 90 % confidence interval. Recovery reduces to fitting a candidate ride against the family, which converts a smartphone into a per-ride vertical-displacement sensor that needs no barometer at inference time, no Wi-Fi survey, and no per-building calibration.

Our pipeline has two stages. Segmentation extracts ride intervals from a continuous, possibly hours-long, inertial stream. Prediction returns Δh for each interval together with a per-ride uncertainty. Both stages read the phone’s accelerometer alone. The per-ride confidence interval is distribution-free, carries finite-sample coverage guarantees, and remains valid under both white sensor noise and the coloured (drift- and vibration-driven) noise that dominates real hand-held recordings.

We evaluate on [TODO: 150] experiments spanning [TODO: TBD] elevators across [TODO: TBD] buildings (residential, office, hotel), with a strict held-out blind test. The system retains its nominal 90 % coverage on the blind test. The median CI half-width stays below [TODO: ± 0.9 m] on short rides, and per-bin error improves on the dataset’s prior best by [TODO: TBD %]. Section 5 reports the full coverage and error statistics.

Contributions.

- (1) **Per-ride displacement with calibrated coverage.** We deliver 90 %, distribution-free, finite-sample confidence intervals on Δh from accelerometer data alone, valid under both white sensor noise and the coloured noise that dominates hand-held recordings.
- (2) **Family-of-basis-functions framework.** We characterise the family of expected gravity-removed acceleration traces for an elevator ride from the cabin’s mechanical limits, and use the same family as the substrate for both segmentation and per-ride prediction.
- (3) **Dataset and harness.** We release the multi-building dataset ([TODO: 150] experiments, multiple phones per ride, barometer ground-truth pipeline), the full evaluation harness, and the operator UI used to triage rides.

Roadmap. Section 2 characterises the elevator’s mechanical structure, the physical limits it imposes on a ride, and the family of expected acceleration functions those limits force. Section 4 presents the two-stage algorithm: it segments an inertial stream into rides and predicts each ride’s Δh with a calibrated interval. Section 3 describes the dataset and its ground truth, and Section 5 reports the train/test split and the results. Sections 6 and 7 discuss limitations and conclude.

2 Background: Elevator Kinematics

This section derives the family of acceleration functions an elevator ride can produce on a smartphone. We start from the cabin's mechanical parts and the physical constraints they impose (§2.1), translate those constraints into bounds on jerk, acceleration, and velocity (§2.2), and then write down the seven-segment S-curve that the bounds force every ride to follow (§2.3). We close with the closed-form relations between machine parameters and the time constants of the ride (§2.4) and the parametric family that this paper's segmentation and prediction stages fit (§2.5); the heavy algebra lives in Appendix A.

2.1 How a passenger elevator works

A passenger elevator is, mechanically, a single moving piece: a cabin (the *car*) hanging on a steel rope that runs up over a sheave at the top of the shaft and down the other side to a counterweight (Figure 1). A motor turns the sheave; the cabin and counterweight rise and fall in opposite directions. The counterweight is sized to match the empty car plus roughly half the rated payload, so the motor only has to drive the small mass difference and the inertia of the moving parts. A brake clamps the sheave shaft whenever the motor is unpowered or anything goes wrong. Three machine pieces decide how the cabin can move: the rope (it stretches), the motor (it has a torque ceiling), and the brake (it has a maximum stopping distance fixed by the headroom h_1). The next subsection turns each into a hard limit on the cabin's motion.




Fig. 1. A traction elevator at a glance. A sheave of radius R couples the car ($m_c + \Delta m_c$) to a counterweight (m_w). The motor delivers torque T_j at angular velocity ω ; rope tensions T_v and T_w load the sheave. H is the hoistway height; h_0 and h_2 are the cabin and counterweight heights; h_1 is the headroom budget above the sheave that the brake must respect.

2.2 Physical limitations and what they bound

Four mechanical realities turn into four hard limits on the cabin's motion. Each is the source of one of the parameters our pipeline fits.

The rope sets the maximum jerk. The main rope stretches. If the motor changes torque too quickly, the cabin starts to bounce on the rope: the passenger feels it as a thump, and the rope itself fatigues over time. Drive controllers therefore ramp torque through a slew-rate limit [10]. Slew-rate-limiting the torque is the same as putting a ceiling on the jerk $\ddot{h}$ the cabin feels. The ride-comfort standard ISO 18738 pins this at $j_{\max} \approx 1.0\text{--}1.6 \text{ m/s}^3$ for passenger elevators [6].

The motor sets the maximum acceleration. At constant velocity the motor only has to fight friction and the small mass mismatch between cabin and counterweight. When it accelerates, it has to push the entire moving mass. The motor has a torque ceiling, so cabin acceleration is also capped: $a_{\max} \approx 0.8\text{--}1.5 \text{ m/s}^2$ on commercial traction elevators. Older geared machines sit at the low end; newer gearless ones at the high end.

The brake sets the maximum cruise velocity. If anything goes wrong — a slack rope, an overspeed sensor trip, loss of power — the cabin has to stop before it hits the top of the shaft. The brake decelerates the cabin at some rate a_{brake} and has at most the headroom h_1 above the top floor to do it in. Energy says $v_{\max}^2 \leq 2 a_{\text{brake}} h_1$, which in practice limits cruise

velocity to $v_{\max} \approx 1.0\text{--}2.5$ m/s in low- and mid-rise buildings (CIBSE Guide D [4]). Faster cabins exist in skyscrapers, but they require taller headroom and shaft-mounted buffers.

Only the height changes from ride to ride. Once the elevator is built, $j_{\max}$, $a_{\max}$, and $v_{\max}$ are fixed for the building — often fixed for the manufacturer and service class. The only thing that changes ride to ride is the floor-to-floor height the cabin has to cover, which we write H (metres, signed for direction). Putting the three limits together, the cabin obeys

$$|\ddot{h}(t)| \leq j_{\max}, \quad |\dot{h}(t)| \leq a_{\max}, \quad |h(t)| \leq v_{\max}, \quad (1)$$

where $\ddot{h} = a$ is the gravity-removed vertical acceleration the phone measures.

2.3 Why the ride is an S-curve

The optimisation problem. The drive solves a simple problem: move the cabin H metres, starting and ending at rest, in the shortest possible time, without ever breaking the three limits in (1). Write the cabin's altitude as $h(t)$, its velocity as $v(t) = \dot{h}(t)$, its acceleration as $a(t) = \ddot{h}(t)$, and the controller's command as the jerk $j(t) = \ddot{h}(t)$. The controller picks $j(t)$; everything else follows by integration. Formally we solve

$$\begin{aligned} & \underset{T, j(\cdot)}{\text{minimize}} && T \\ & \text{subject to} && h(0) = 0, \quad h(T) = H, \quad v(0) = v(T) = 0, \\ & && |j(t)| \leq j_{\max}, \quad |a(t)| \leq a_{\max}, \quad |v(t)| \leq v_{\max}, \\ & && \forall t \in [0, T]. \end{aligned} \quad (2)$$

The result. Pontryagin's maximum principle, applied to (2), forces the optimal command to be *bang-bang* on the highest-derivative input. At every instant, the optimal jerk takes only three values,

$$j^*(t) \in \{-j_{\max}, 0, +j_{\max}\}, \quad (3)$$

with $j^* = 0$ holding the cabin against $|a| = a_{\max}$ or $|v| = v_{\max}$ when one of the lower-order limits is itself saturated. Intuitively, any moment spent at an interior value $|j| < j_{\max}$ is wasted: the cabin could have been pushing harder and finished sooner. Appendix A.1 sketches the derivation; we use (3) below.

The S-curve and its three regimes. Stitching the bang-bang jerk command (3) together with the three limits yields the canonical *S-curve* acceleration profile. The floor-to-floor height H alone fixes how many of the limits a ride saturates, and so splits every ride into one of three regimes — long, medium, or short — shown side by side in Figure 2.


2.4 From machine parameters to ride-time constants

The bang-bang structure pins every phase of the ride to a closed-form duration in the machine parameters and the floor-to-floor height H . Three durations describe any ride:

$$T_j = \frac{a_{\max}}{j_{\max}}, \quad (4)$$

$$T_a = \max\left(0, \frac{v_{\max}}{a_{\max}} - T_j\right), \quad (5)$$

$$T_v = \max\left(0, \frac{H}{v_{\max}} - 2T_j - T_a\right), \quad (6)$$

**Long ride — 7 segments.**

$$T_v > 0, \quad T_a > 0$$

All three limits saturate; a flat $v_{\max}$ cruise splits the two trapezoidal lobes.

Medium ride — 5 segments.

$$T_v = 0, \quad T_a > 0$$

Jerk and acceleration saturate but the cabin never reaches $v_{\max}$; the two lobes touch at the peak-velocity instant.

Short ride — 3 segments.

$$T_v = T_a = 0$$

Only jerk saturates; the two lobes are triangles that meet, with peak $|a| < a_{\max}$.

Fig. 2. The three S-curve regimes of the time-optimal acceleration profile $a(t)$, all produced by the bang-bang jerk command (3) and selected by the floor-to-floor height H . Each row pairs the profile with the time-constant condition that defines its regime; longer H saturates more limits. Closed forms for T_j , T_a , and T_v follow in §2.4.

and they sum over the seven phases to the total ride duration

$$T_{\text{ride}} = 4T_j + 2T_a + T_v = \frac{H}{v_{\max}} + \frac{v_{\max}}{a_{\max}} + \frac{a_{\max}}{j_{\max}} \quad (\text{trapezoidal regime}). \quad (7)$$

Figure 3 locates the four durations on a single ride. The jerk ramp T_j and the acceleration plateau T_a are constants of the building. T_j depends only on the machine — about 0.77 s for a typical lift ($a_{\max} = 1.0 \text{ m/s}^2$, $j_{\max} = 1.3 \text{ m/s}^3$) — and T_a vanishes on rides too short to reach $v_{\max}$. Only the cruise T_v scales with ride length, growing as $H/v_{\max}$. Equation (7) is the inversion handle later sections use to recover H jointly with the machine parameters; Appendix A derives all four constants from the bounds in (1).

2.5 The trapezoid basis function and the parametric family

Combining (1)–(7), every ride is fixed by four parameters,

$$\theta = (j_{\max}, a_{\max}, v_{\max}, H), \quad (8)$$

three constants of the building and the ride-specific height H .




Fig. 3. The four ride-time constants on the acceleration trace $a(t)$ of a trapezoidal ride: jerk ramp T_j , the $a_{\max}$ plateau T_a , the $v_{\max}$ cruise T_v (the flat $a = 0$ gap between lobes), and the total ride T_{ride} .

The trapezoid basis function. Each acceleration lobe of the S-curve is a trapezoid in time — a ramp up, a flat top, and a ramp down. We normalise it to the trapezoid basis function

$$\tau_{W,f}(t) = \begin{cases} 1 & |t| \leq fW, \\ \frac{W - |t|}{(1-f)W} & fW < |t| \leq W, \\ 0 & |t| > W, \end{cases} \quad (9)$$

of half-width W and flat fraction $f \in [0, 1]$: a unit plateau over $[-fW, fW]$ joined to zero by two linear ramps (Figure 4).

Mapping to the machine parameters. The first lobe spans $2T_j + T_a$, so the machine parameters fix (W, f) exactly:

$$W = T_j + \frac{1}{2}T_a = \frac{a_{\max}}{2j_{\max}} + \frac{1}{2}\max\left(0, \frac{v_{\max}}{a_{\max}} - T_j\right), \quad (10)$$

$$f = \frac{T_a}{T_a + 2T_j} = \frac{\max(0, v_{\max}/a_{\max} - T_j)}{\max(0, v_{\max}/a_{\max} - T_j) + 2T_j}. \quad (11)$$

A short hop that never reaches $v_{\max}$ has $T_a = 0$ and $f = 0$, a pure triangle; a long ride has $T_a \gg T_j$ and $f \rightarrow 1$, a near-rectangular plateau. The two lobes pull apart — their centroid spacing Δt_c exceeds $2W$ — only once a $v_{\max}$ cruise opens between them, the 7-segment regime of §2.3.

The parametric family. Every ride is therefore a pair of opposite-sign trapezoids of shared shape,

$$a_\theta(t) = sA\tau_{W,f}(t - t_{c,1}) - sA\tau_{W,f}(t - t_{c,2}), \quad (12)$$

with amplitude $A = a_{\max}$, direction $s \in \{+1, -1\}$, and centroid spacing $\Delta t_c = t_{c,2} - t_{c,1} = H/v_{\max}$ (Appendix A).

Three facts about this family $\mathcal{F}_{\text{ride}}$ drive the rest of the paper.

(i) Three of its parameters are constants of the building, so a ride differs from the building's other rides only in the height H — the single quantity the predictor of §4 recovers. (ii) The two lobes carry equal and opposite area, so $\int_0^{T_{\text{ride}}} a_\theta(t) dt = 0$ — the zero-mean anchor for any ZUPT-style predictor. (iii) The lobe spacing $\Delta t_c = H/v_{\max}$ is a scale-free estimator of H , so the matched-filter detector needs no per-building amplitude calibration.




Fig. 4. The trapezoid basis function $\tau_{W,f}(t)$ of (9).

2.6 The smartphone accelerometer

The pipeline reads a single phone sensor — the accelerometer — and one obstacle separates its raw signal from the vertical acceleration a ride produces: the phone's unknown orientation.

The accelerometer. Every phone carries a six-axis MEMS IMU [3]. Its TYPE_LINEAR_ACCELERATION stream reports the gravity-removed acceleration as a three-vector $\mathbf{a} = (a_x, a_y, a_z)$ along the phone's *body axes* — the x, y, z fixed to the case (Figure 5). The sensor-hub rate varies by device, so we resample every recording to a uniform 100 Hz grid ($\Delta t = 10$ ms). The noise floor is white, but a hand-held phone adds coloured drift and 1–50 Hz vibration, so later sections fit the noise model on real ride residuals.




Fig. 5. Body axes x, y, z tilt from the world vertical $\hat{z}_w$ by an unknown rotation R .

The unknown vertical. An elevator ride is purely vertical: the cabin accelerates along the world up-axis $\hat{z}_w$, and that one scalar $a(t)$ is all the pipeline wants. The accelerometer, though, measures along the phone’s body axes, and a hand-held phone tilts arbitrarily. An unknown rotation R separates the two frames, so the vertical signal spreads across all three channels, $\mathbf{a} = R^T a \hat{z}_w$ — the body channel a_z is not the world vertical. Recovering $a(t)$ means recovering which body direction points up; the accelerometer-only pipeline reads it from gravity in the stationary windows around each ride.

In short, the pipeline works from a single 100 Hz body-frame acceleration stream.

3 Dataset and Ground Truth

The dataset is a field campaign of [TODO: ≈ 150] elevator rides recorded across [TODO: 7] elevators in [TODO: 6] buildings in central Israel (Table 2). The buildings span residential, office, and hotel cabins, each with a different cabin geometry, motor, and floor-to-floor height, so no estimator can overfit a single machine. Four phones, carried together, recorded every ride, so one physical event produced four heterogeneous accelerometer traces — different MEMS parts, sampling jitter, and noise floors. At least one of the four phones carries a barometric pressure sensor, which serves as the ground-truth device even when the phone under test has none.

Every ride’s signed displacement Δh comes from the barometer rather than from hand-labelling. We invert the pressure stream to altitude through the ISA barometric formula; a consumer barometer resolves altitude to roughly 10 cm, an order of magnitude finer than the ~ 3 m inter-floor pitch the pipeline must separate. Low-pass filtering and differencing the same altitude trace yields the up, down, and outside interval labels that cut each session into discrete rides. A phone without a barometer borrows the pressure stream of a co-carried device, aligned to it by accelerometer cross-correlation (Appendix B).

Raw barometric altitude drifts with weather and building airflow, so we snap it to the building itself. Each building contributes a *gramushka* table — the floor elevations read from its accordion-folded architectural drawing (Figure 6). We anchor every session at its known start floor and integrate the barometer altitude across each detected ride. Snapping the ride endpoint to the nearest table floor makes the ground-truth Δh a true floor-to-floor displacement. The pipeline flags any snap whose distance exceeds [TODO: TBD] m for manual review rather than forcing it, and the one elevator without a drawing falls back to the drift-corrected barometer Δh (Appendix B).

4 Algorithm

The pipeline converts a continuous, hours-long accelerometer stream into one signed displacement Δh per elevator ride, each carrying a calibrated 90 % confidence interval, in two stages (Figure 7). Segmentation slices the stream into discrete ride intervals (§4.1); prediction fits each interval for its Δh and uncertainty (§4.2). Both stages reuse the trapezoid family of §2.5 directly: a ride is an opposite-sign pair of trapezoid lobes (12), so the matched filter that finds a ride is the matched filter that measures it.




Fig. 6. Snapping the drifting barometer altitude (blue) to the nearest *gramushka* floor line (dashed) yields a drift-free ground-truth Δh .




Fig. 7. The two-stage pipeline. Segmentation turns a continuous accelerometer stream into discrete ride intervals; prediction turns each interval into a signed displacement Δh with a calibrated confidence interval. Both stages fit the trapezoid ride family of §2.5.

4.1 Segmentation

Segmentation cuts an hours-long accelerometer stream into discrete ride intervals, each tagged up or down. A ride is an opposite-sign pair of trapezoid lobes (12), so segmentation detects candidate lobes (§4.1.1) and then pairs them into rides (§4.1.2).

4.1.1 Detecting candidate lobes. The pipeline projects the three-axis accelerometer onto the session’s gravity direction (§2.6) and smooths it into a signed vertical trace $a(t)$. A least-squares matched filter scores $a(t)$ against every lobe template $\tau_{W,f}$ of the family $\mathcal{F}_{\text{ride}}$ (§2.5): at each sample i , fitting one amplitude to the length- $2W$ window $a_{[i]}$ centred on i gives

$$\hat{A}_{W,f}(i) = \frac{\langle a_{[i]}, \tau_{W,f} \rangle}{\|\tau_{W,f}\|^2}, \quad R_{W,f}^2(i) = 1 - \frac{\|a_{[i]} - \hat{A}_{W,f}(i) \tau_{W,f}\|^2}{\|a_{[i]}\|^2}. \quad (13)$$

Collapsing the (W, f) grid by amplitude sign yields, per sample,

$$R^2(i) = \max_{W,f} R_{W,f}^2(i), \quad R_{\pm}^2(i) = \max_{\pm \hat{A}_{W,f}(i) > 0} R_{W,f}^2(i), \quad (14)$$

where a high $R_+^2(i)$ marks a take-off lobe and a high $R_-^2(i)$ a landing lobe. A candidate lobe is a strict local maximum of R_+^2 or R_-^2 that clears an amplitude and a shape floor,

$$|\hat{A}(i)| \geq a_{\min}, \quad R^2(i) \geq r_{\min}, \quad (15)$$

with $\hat{A}(i)$ the amplitude at the best grid cell; non-maximum suppression then drops near-duplicates within δ_{nms} and keeps same-sign lobes at least Δ_{ss} apart (Figure 8). The grid bounds (10)–(11) and the thresholds are calibrated on the training set (Appendix C).




Fig. 8. Signed score traces $R_+^2(t)$ (take-off, blue) and $R_-^2(t)$ (landing, red) around one ride. Local maxima are candidate lobes; dot colour marks the first stage to reject each peak, leaving the committed take-off/landing pair (green).

4.1.2 *Pairing candidate lobes into rides.* Each candidate lobe i admits an opposite-sign partner set inside the ride-duration window,

$$\mathcal{P}(i) = \{ j : \text{sign } \hat{A}(j) = -\text{sign } \hat{A}(i), \quad t_{\min} \leq |t_j - t_i| \leq t_{\max} \}. \quad (16)$$

A genuine ride's two lobes are cut from one cabin profile, so each pair (i, j) is fitted with a single template and a single shared amplitude. With a_k the k th lobe window, $P_k = \|a_k\|^2$, $u_k = s_k \langle a_k, \tau_{W,f} \rangle$ its sign-corrected correlation, and $N = \|\tau_{W,f}\|^2$,

$$A^*(W, f) = \frac{u_1 + u_2}{2N}, \quad R_k^2(W, f) = 1 - \frac{P_k - 2A^*u_k + (A^*)^2N}{P_k}, \quad (17)$$

and the pair score is $S(i, j) = \max_{W,f} \frac{1}{2}(R_1^2 + R_2^2)$, attained at (W^*, f^*, A^*) . These starred values describe one *average* lobe: the fit assumes take-off and landing share a shape, so A^* averages the two single-lobe amplitudes. A high S means both lobes trace out that same trapezoid. The pair is admissible when

$$S(i, j) \geq r_{\text{pair}}, \quad A^* \geq a_{\text{pair}}, \quad E(i, j) \geq e_{\min}, \quad \text{RMS}_{\text{mid}} \leq \rho A^*. \quad (18)$$

The grid energy E averages the clamped joint R^2 over every cell of the (W, f) template grid G ,

$$E(i, j) = \frac{1}{|G|} \sum_{(W,f) \in G} \max(0, \frac{1}{2}(R_1^2(W, f) + R_2^2(W, f))), \quad (19)$$

so a real ride — which lights a broad band of cells, not one — clears the floor $e_{\min}$. The cruise term RMS_{mid} is the trace RMS on the inter-lobe stretch, near zero while the cabin holds constant velocity. A greedy resolver ranks admissible pairs by

$$\text{rank}(i, j) = S(i, j) - \lambda |t_j - t_i| \quad (20)$$

and commits them highest-first, skipping any that reuse a lobe or overlap a committed interval. The penalty λ defeats the *super pair* — a take-off matched to a far-later landing, which scores like a true ride but spans far longer. The committed pairs are the ride intervals passed to prediction.

Appendix C explains and illustrates every detection and pair-filter threshold.

4.2 Prediction

Prediction takes one ride interval and returns the signed displacement Δh , a per-ride standard deviation σ , a calibrated 90 % confidence interval, and an accepted flag. Displacement follows from one fact: an accelerometer measures acceleration, and displacement is its double integral — integrate the vertical trace $a(t)$ once for velocity, twice for Δh . The front end recovers that signed vertical trace from the three-axis accelerometer by the gravity projection of §2.6.

Two accelerometer-only algorithms run that double integral by different routes. ZUPT double integration (§4.2.1) integrates the measured $a(t)$ directly, anchored by the cabin resting at both ends of the ride. The trapezoid pulse-pair fit (§4.2.2) integrates a model instead: it constructs the ride's theoretical acceleration signal — the trapezoid pulse pair of §2.5 — and integrates it in closed form, reading Δh off a formula in the fitted shape. Section 5 compares them head to head.

Both feed three shared stages: a per-ride error model (§4.2.3), a conformal layer that calibrates σ into the interval (§4.2.4), and a quality filter that sets the accepted flag (§4.2.5). The trapezoid fit borrows only ZUPT's cruise-velocity scale as an amplitude anchor, never its displacement, so the two stay independent estimators and the comparison of §5 stays meaningful.

4.2.1 *ZUPT double integration.* ZUPT recovers Δh by double-integrating the measured vertical acceleration under a zero-velocity boundary at both ends of the ride,

$$v(t) = \int_0^t a(\tau) d\tau - \frac{t}{T} \int_0^T a(\tau) d\tau, \quad \Delta h = \int_0^T v(t) dt, \quad (21)$$

over a ride of duration T .

The boundary assumption anchors the integral. A ride begins and ends with the cabin at rest, so $v(0) = v(T) = 0$. The first integral of (21) enforces $v(0) = 0$ by construction; the linear ramp term subtracts the integrated-noise drift that would otherwise leave $v(T) \neq 0$ — the zero-velocity update that names the method.

ZUPT integrates only the moving span: the pipeline thresholds the vertical trace at $|a(t)| = 0.10 \text{ m/s}^2$ and zeroes the quiet pre- and post-roll, so noise outside the ride never enters the integral. The count N of active samples sets the scale of the error model (§4.2.3).

4.2.2 *Trapezoid pulse-pair fit.* The trapezoid fit recovers Δh from the *shape* of the acceleration trace. The ride’s theoretical signal is the S-curve of §2.3, set by three machine parameters ($j_{\max}$, $a_{\max}$, $v_{\max}$) besides the height. Fitting all three to a hand-held recording is ill-conditioned: the jerk ramps occupy too small a fraction of the ride to constrain them. The fit instead collapses each acceleration lobe to the trapezoid $\tau_{W,f}$ the detector already uses, a two-parameter shape the bang-bang jerk command makes exact (§2.5).

The segmenter (§4.1.2) already returns the fitted parameters ($W, f, A, s, t_{c,1}, t_{c,2}$) that reconstruct the ride’s acceleration as the opposite-sign trapezoid pair

$$a_\theta(t) = sA \tau_{W,f}(t - t_{c,1}) - sA \tau_{W,f}(t - t_{c,2}), \quad (22)$$

the estimated ride signal (Figure 9). Double-integrating $a_\theta(t)$ gives Δh in closed form,

$$\Delta h = sAW(1+f)\Delta t_c, \quad \Delta t_c = t_{c,2} - t_{c,1}; \quad (23)$$

Appendix D works through the integral and the joined-pulse limit.

4.2.3 *Theoretical error analysis.* Every prediction carries a per-ride standard deviation σ derived from its own fit, not a constant. The two algorithms reach that σ from different error sources.

ZUPT. Double integration accumulates sensor noise. Integrating white noise of standard deviation σ_a twice over N active samples of spacing Δt gives an end-point displacement variance whose root is

$$\sigma_{\text{white}} = \sigma_a \Delta t^2 \sqrt{N^3/12}, \quad (24)$$

the classic ZUPT result, with the factor $1/12$ crediting the linear drift correction of (21) for removing one degree of freedom. A hand-held phone carries three further error sources, which the total budget adds to the white-noise term in quadrature,

$$\sigma_{\text{ZUPT}}^2 = \sigma_{\text{white}}^2 + \sigma_{\text{mech}}^2 + \sigma_{\text{rel}}^2 + \sigma_{\text{proj}}^2 : \quad (25)$$




Fig. 9. The reconstructed signal $a_\theta(t)$ (red) overlaid on the smoothed vertical trace (black) for one ride: two opposite-sign trapezoid lobes centred at $t_{c,1}$ and $t_{c,2}$.

a mechanical-jitter term σ_{mech} for hand and pocket movement, a relative-scale term σ_{rel} that grows with $|\Delta h|$, and a projection penalty σ_{proj} for rides that lack a clean gravity reference. Appendix D derives all four.

Trapezoid. The trapezoid budget propagates the matched-filter fit uncertainty through the closed-form height (23) by the delta method,

$$\sigma_{\Delta h}^2 = [W(1+f) \Delta t_c]^2 \sigma_A^2 + [A(1+f) \Delta t_c]^2 \sigma_W^2 + [AW \Delta t_c]^2 \sigma_f^2 + [AW(1+f)]^2 \sigma_{\Delta t_c}^2, \quad (26)$$

with the per-parameter variances supplied by the matched-filter Cramér–Rao bound (template energy and noise level). Two data-adaptive corrections fit the budget to hand-held data: the effective noise level uses the fit residual rather than the datasheet figure, and the whole budget is divided by the joint R^2 . Appendix D derives every term.

4.2.4 Confidence interval. The per-ride σ sets the *shape* of the uncertainty but not its *scale*: nothing so far guarantees that $\pm\sigma$ brackets the true height with any stated probability. Split-conformal calibration [8] supplies the guarantee. On a held-out calibration set the pipeline computes, for every ride i , the nonconformity score

$$r_i = \frac{|\widehat{\Delta h}_i - \Delta h_i|}{\sigma_i}, \quad (27)$$

the prediction error in units of its own σ , and takes the empirical $(1 - \alpha)$ quantile of $\{r_i\}$ as a single multiplier k . The reported interval is $\widehat{\Delta h} \pm k\sigma$. By the standard split-conformal argument the interval has distribution-free, finite-sample coverage,

$$\mathbb{P}(|\widehat{\Delta h} - \Delta h| \leq k\sigma) \geq 1 - \alpha, \quad (28)$$

for exchangeable rides, with no assumption on the error distribution. We set $\alpha = 0.1$ for a 90 % interval.

The same machinery wraps both algorithms, but each calibrates its multiplier separately. ZUPT calibrates k on the clean calibration rides; the trapezoid fit calibrates on the rides that are both clean and quality-accepted, its exchangeable population. The multiplier is the $\lceil (1 - \alpha)(n + 1) \rceil / n$ empirical quantile of the n scores, floored at 1.645 so a small calibration sample cannot return an interval narrower than the Gaussian 90 % half-width.

The interval is *adaptive* because the multiplier is shared while σ is not. A long, cleanly fitted ride and a short, poorly fitted one take the same k but very different σ , so the interval widens exactly where the error model says the fit is weak — short rides, unstable gravity, low joint R^2 . Appendix D states the exchangeability assumption and the per-algorithm calibration protocol.

4.2.5 Quality filter. The last step attaches the accepted flag. Each algorithm runs its own quality filter, but the two share a shape: a handful of interpretable features, a set of hard gates, and an additive penalty score over the graded features. The first breach rejects the ride with a human-readable reason, and the filter rejects whenever any hard gate trips or the penalty score exceeds 6.

The two filters differ in what they inspect. Both check the bracketing-window gravity, the phone’s orientation, and the active-motion fraction — any accelerometer estimate is only as good as its vertical axis. ZUPT adds an impact gate and an end-velocity check (Table 4). The trapezoid filter adds fit-quality and geometry checks: the joint R^2 , the residual autocorrelation, the agreement between the velocity-anchored and matched-filter amplitudes, the residual energy outside the two lobes, and the steadiness of the integrated cruise velocity (Table 5). Appendix D explains and illustrates every check.

The quality filter is a *filter, not a predictor*. It targets the rides where the σ model must under-predict the error — a phone rotated mid-ride, an impact that swamps the lobes, a fit with no pulse structure. On those rides no interval width

is trustworthy, so the pipeline declines to report rather than report a confident wrong answer. Everywhere else the conformal interval of §4.2.4 carries the coverage guarantee, so the filter passes almost every clean ride and rejects only the pathological.

5 Evaluation

6 Discussion

7 Conclusion

References

- [1] Michael Athans and Peter L. Falb. 1966. *Optimal Control: An Introduction to the Theory and Its Applications*. McGraw-Hill.
- [2] Paramvir Bahl and Venkata N. Padmanabhan. 2000. RADAR: An In-Building RF-Based User Location and Tracking System. In *Proceedings of IEEE INFOCOM*. 775–784. doi:10.1109/INFCOM.2000.832252
- [3] Stephen Beeby, Graham Ensell, Michael Kraft, and Neil White. 2004. *MEMS Mechanical Sensors*. Artech House.
- [4] Chartered Institution of Building Services Engineers. 2020. *CIBSE Guide D: Transportation Systems in Buildings* (fifth ed.). CIBSE.
- [5] Robert Harle. 2013. A Survey of Indoor Inertial Positioning Systems for Pedestrians. *IEEE Communications Surveys & Tutorials* 15, 3 (2013), 1281–1293. doi:10.1109/SURV.2012.121912.00075
- [6] International Organization for Standardization. 2012. ISO 18738-1:2012 — Measurement of Ride Quality, Part 1: Lifts (Elevators). Standard, ISO.
- [7] Donald E. Kirk. 2004. *Optimal Control Theory: An Introduction* (reprint ed.). Dover Publications.
- [8] Jing Lei and Larry Wasserman. 2014. Distribution-Free Prediction Bands for Non-Parametric Regression. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 76, 1 (2014), 71–96. doi:10.1111/rssb.12021
- [9] Hui Liu, Houshang Darabi, Pat Banerjee, and Jing Liu. 2007. Survey of Wireless Indoor Positioning Techniques and Systems. *IEEE Transactions on Systems, Man, and Cybernetics, Part C* 37, 6 (2007), 1067–1080. doi:10.1109/TSMCC.2007.905750
- [10] Sonja Macfarlane and Elizabeth A. Croft. 2003. Jerk-Bounded Manipulator Trajectory Planning: Design for Real-Time Applications. *IEEE Transactions on Robotics and Automation* 19, 1 (2003), 42–52. doi:10.1109/TRA.2002.807548
- [11] Kartik Muralidharan, Azeem Javed Khan, Archan Misra, Rajesh Krishna Balan, and Sharad Agarwal. 2014. Barometric Phone Sensors: More Hype Than Hope!. In *Proceedings of the 15th Workshop on Mobile Computing Systems and Applications (HotMobile)*. 12:1–12:6. doi:10.1145/2565585.2565596
- [12] George R. Strakosch and Robert S. Caporale. 2010. *The Vertical Transportation Handbook* (fourth ed.). Wiley.

A Theoretical Elevator Pulse Analysis

This appendix derives the constants of §2.4 from the bounds in (1). §A.1 shows the optimal jerk is bang-bang; §A.2 integrates the resulting seven-segment profile to obtain T_j , T_a , T_v , T_{ride} , and the centroid-spacing identity used in §2.5.

A.1 The optimal elevator pulse

The time-optimal control problem (2) forces the optimal jerk to take only three values, $j^*(t) \in \{-j_{\max}, 0, +j_{\max}\}$ (3). We sketch the three moves Pontryagin’s maximum principle turns on; the full machinery is standard for time-optimal control of a chain of integrators with bounded input [1, 7].

Hamiltonian. Stack the state $x = (h, v, a)^\top$ with dynamics $\dot{x} = (v, a, j)^\top$, control j , free terminal time T , and unit running cost. Adjoining co-states $\lambda = (\lambda_h, \lambda_v, \lambda_a)^\top$ gives

$$\mathcal{H}(x, \lambda, j) = 1 + \lambda_h v + \lambda_v a + \lambda_a j,$$

which is linear in j .

Switching law. Linearity in j pins the PMP-minimising control to the boundary $|j| = j_{\max}$ whenever $\lambda_a \neq 0$:


$$j^*(t) = -j_{\max} \operatorname{sign} \lambda_a(t).$$

The co-state equations $\dot{\lambda} = -\partial H / \partial x$ yield $\lambda_a(t) = \frac{1}{2} \lambda_h t^2 - \lambda_v(0) t + \lambda_a(0)$, a quadratic that crosses zero at most twice; the unconstrained j^* therefore switches sign at most twice.

Constrained arcs. When $|a| = a_{\max}$ activates, the cabin rides the boundary and the controller must hold $\dot{a} = j = 0$; the velocity constraint $|v| = v_{\max}$ similarly forces $a = 0$ and hence $j = 0$. Stitching the unconstrained $\pm j_{\max}$ arcs with these constrained zero-jerk arcs yields (3) and the three regimes selected by the floor-to-floor height H (Figure 2).

A.2 Time parameters formula development

We integrate the bang-bang jerk over the take-off lobe (phases 1–3), recover the three durations, and sum to the total ride time. We assume the trapezoidal regime $T_v > 0$; the medium and short regimes are limits $H \downarrow a_{\max} T_j (T_j + T_a)$ of the same formulas. Each phase paints one slab of the lobe beside: the two red wedges of phases 1 and 3, the orange plateau of phase 2. Their areas are the velocity gained on that phase, and the colours below match the equations.



Phase 1 — jerk ramp ($j = +j_{\max}$). With $a(0) = v(0) = 0$, one integration gives

$$a(t) = j_{\max} t, \quad a(T_j) = a_{\max} \implies T_j = \frac{a_{\max}}{j_{\max}} \quad (4).$$

Integrating twice more,

$$v(T_j) = \frac{1}{2} a_{\max} T_j, \quad h(T_j) = \frac{1}{6} a_{\max} T_j^2.$$

The velocity $v(T_j)$ is the area of the phase-1 red wedge.

Phase 2 — acceleration plateau ($a = a_{\max}$). On $t \in [T_j, T_j + T_a]$ the cabin holds at $a_{\max}$, so the velocity gains the area of the orange rectangle,

$$v(T_j + T_a) = v(T_j) + a_{\max} T_a = a_{\max} \left(\frac{1}{2} T_j + T_a \right),$$

and the displacement added is $\Delta h_2 = v(T_j) T_a + \frac{1}{2} a_{\max} T_a^2$.

Phase 3 — jerk ramp-down ($j = -j_{\max}$). On $t \in [T_j + T_a, 2T_j + T_a]$, $a(t) = a_{\max} - j_{\max}(t - T_j - T_a)$ so $a(2T_j + T_a) = 0$ as required. Phase 3 contributes a second $\frac{1}{2} a_{\max} T_j$ of velocity — the mirror red wedge — so the lobe-end velocity is the sum of three slab areas,

$$v(2T_j + T_a) = \frac{1}{2} a_{\max} T_j + a_{\max} T_a + \frac{1}{2} a_{\max} T_j = a_{\max} (T_j + T_a).$$

Demanding the cruise enter at the velocity ceiling, $a_{\max} (T_j + T_a) = v_{\max}$, gives $T_a = \frac{v_{\max}}{a_{\max}} - T_j$ (5). The displacement added in phase 3 is $\Delta h_3 = v(T_j + T_a) T_j + \frac{1}{2} a_{\max} T_j^2 - \frac{1}{6} j_{\max} T_j^3$.

Synthesis. Summing phases 1–3 with $j_{\max} T_j = a_{\max}$ and $a_{\max} (T_j + T_a) = v_{\max}$ collapses to

$$d_{\text{accel}} = h(T_j) + \Delta h_2 + \Delta h_3 = \frac{1}{2} v_{\max} (2T_j + T_a).$$

By symmetry phases 5–7 contribute the same d_{accel} . The cruise covers the remainder $H - 2d_{\text{accel}}$ at $v_{\max}$,

$$T_v = \frac{H - v_{\max} (2T_j + T_a)}{v_{\max}} = \frac{H}{v_{\max}} - 2T_j - T_a \quad (6),$$

and $T_{\text{ride}} = 4T_j + 2T_a + T_v$ collapses to the additive form (7).

Centroid spacing of the two acceleration lobes. The first lobe spans $[0, 2T_j + T_a]$ with centroid at $t_{c,1} = T_j + T_a/2$ (trapezoid symmetry); the second spans $[2T_j + T_a + T_v, 4T_j + 2T_a + T_v]$ with centroid at $t_{c,2} = 3T_j + 3T_a/2 + T_v$. Subtracting and using (6),

$$\Delta t_c = t_{c,2} - t_{c,1} = 2T_j + T_a + T_v = \frac{H}{v_{\max}}.$$

The temporal gap between the two acceleration lobes equals $H/v_{\max}$ and is independent of $j_{\max}$ and $a_{\max}$ — the amplitude-free centroid identity used in the basis-function mapping of §2.5.

B Dataset collection and ground-truth pipeline

The dataset is 195 labelled elevator rides drawn from eight recording sessions across seven elevators in six buildings, all collected in March–April 2026. This appendix walks the construction pipeline end-to-end: how each session was recorded and how every ride got its label and floor-to-floor Δh (§B.1), the human-in-the-loop edit pass that corrects detector mistakes (§B.2), and the per-session statistics of the campaign (§B.3).

B.1 Method

Recording setup. Each session ran four phones simultaneously, carried together in one bag or hand so every device saw the same motion, vibration, and acoustics. Two experimenters carried two phones each, and the fleet always included at least one barometer-equipped device whose pressure trace became the ground truth for the three accelerometer-only phones. Two of the eight sessions — Haari 3 and Bar-Ilan 2 — relied on a single experimenter holding one barometer-equipped phone. Table 1 lists the channels each phone logged at its native rate; the loader resamples every channel to 50 Hz before downstream stages see it.

Barometric pressure to altitude. The recording phone’s pressure channel converts to a metric altitude trace through the International Standard Atmosphere troposphere inversion, $h = (T_0/L) [1 - (P/P_0)^{1/5.255}]$, with the surface temperature T_0 read from the session log. Figure 10 plots both sides on a Millennium-Hotel session. Pressure climbs by ~ 6 hPa as the experimenter descends from floor 12 to the ground lobby, and the inverted altitude resolves the matching ~ 50 m descent on the same time axis. At standard T_0 the resolution sits near 10 cm per sample, an order of magnitude below the ~ 3 m typical inter-floor pitch, so floor-to-floor steps stand clear of sensor noise.




Fig. 10. Pressure (**top**) and the ISA-inverted altitude (**bottom**) on a Millennium-Hotel session. The pressure rise around $t \approx 1000$ s is the experimenter taking the elevator down from floor 12; the inverted altitude shows the matching ~ 50 m descent.

Edge-point ride extraction. HeightSegmenter turns the altitude trace into ride intervals through a deterministic state machine over vertical velocity. The procedure low-pass filters altitude across an 8 s rolling window, differences the result into velocity, smooths velocity with a 3 s box kernel, and marks samples above $|v_z| > 0.15$ m/s as moving. Same-sign moving runs within 6 s of each other merge into one interval. The detector then drops any survivor shorter than 3 s or with $|\Delta z| < 2$ m. The pipeline pads each remaining survivor by 1 s on either side and labels it up when the mean velocity is positive, down otherwise; gaps carry the outside label, so every sample wears one of the three tags. Figure 11 runs the procedure on a 170 s window of Millennium-Hotel session exp2. Four step-shaped descents in the altitude panel become four negative-velocity dips that breach the threshold, and the bottom panel emits the matching down intervals interleaved with outside gaps. Every operation is deterministic, so repeated runs label a recording identically.




Fig. 11. Barometer interval detector on a 170 s window of Millennium-Hotel session exp2. **Top:** low-pass altitude. **Middle:** smoothed vertical velocity with the ± 0.15 m/s threshold (dashed) and the “moving” regions shaded. **Bottom:** the surviving up/down intervals after merging, with outside between them.

Sensor channels. An Android logger subscribes to every inertial and environmental channel the phone exposes and writes one tab-separated row per sample (Table 1). The pipeline reads accelerometer (ACC) on every phone, barometer (PRS) on the equipped phone for ground truth, and gyroscope, magnetometer, and orientation for the editor’s secondary plots and the cross-device alignment step.

Tag	Sensor	Columns
ACC	Linear accelerometer	x, y, z
GYR	Gyroscope	x, y, z
MAG	Magnetometer	x, y, z
ORI	Rotation-vector quaternion	w, x, y, z
PRS	Barometer	pressure
GPS	GPS fix	lat, lon, alt

Table 1. Sensor channels emitted by the logger. The pipeline reads ACC throughout; PRS supplies ground truth only.

Cross-device time synchronisation. Android stamps each sample against device boot time, not wall-clock, so two phones in the same session disagree on event time by their boot-difference — often tens of thousands of seconds. The loader anchors each log on its first-sample timestamp to remove the gross offset, then sharpens the alignment with a normalised cross-correlation of the two accelerometer magnitudes across a ± 10 s window. After correction the residual lag sits below one 50 Hz sample, and the pipeline bakes the recovered offset into the secondary barometer stream before merging.

Gramushka tables and floor snapping. Raw barometer altitude drifts with weather and building airflow, so the pipeline snaps every ride endpoint to a building-specific floor table. Figure 12 threads the three steps. Each building contributes a *gramushka* — the team’s nickname for the accordion-folded architectural section drawing (Figure 12a). Transcribing the drawing’s labelled floor lines yields a lookup from floor name to surveyed elevation (Figure 12b). For each detected ride the pipeline integrates the temperature-aware altitude across the interval, snaps the endpoint to the nearest floor in the lookup, and treats the resulting floor-to-floor displacement as the ground-truth Δh (Figure 12c; the worked example reads raw +29.29 m and snaps +1.49 m down to Floor 5 at +27.80 m). The external Millennium elevator has no drawing; its rides keep the drift-corrected barometer Δh directly.




Fig. 12. From drawing to floor-height lookup to snapped ride endpoint. (a) The accordion-folded architectural section drawing for one campaign building (Bar-Ilan 2, Herzeliya). (b) The transcribed lookup, mapping each floor name to one surveyed elevation (annotated arrow: Floor 10 at +42.80 m). (c) A worked snap from a real Bar-Ilan 2 ride: the raw barometer altitude +29.29 m snaps +1.49 m down to Floor 5 at +27.80 m.

B.2 Hand editing

Why hand editing. The barometer detector cannot tell elevator motion from any other vertical motion of the same magnitude. Stair climbs, driving up or down a hill, and HVAC pressure ripples each pass the velocity threshold and slip into the label stream; conversely, an over-tight merge fuses two close rides into one. Experimenters carry side information — which button they pressed, which lobby they entered, which staircase they took — that no sensor sees, and the hand-edit pass injects that knowledge into the labels.

The editor. Every recording passes through the ground-truth editor of Figure 13 at least once. Stacked panels show altitude, vertical velocity, and the accelerometer, gyroscope, and magnetometer magnitudes on a shared time axis; ride

intervals overlay as coloured spans (green = up, red = down, grey = outside). A side table lists every interval for direct editing, and the detector’s current proposals appear as toggleable hatched bands the editor can promote, redraw, or reject.



Fig. 13. The ground-truth editor. Stacked panels show altitude, vertical velocity, and the accelerometer, gyroscope, and magnetometer magnitudes on a shared time axis; ground-truth intervals are overlaid as coloured spans (green = up, red = down, grey = outside). The side table lists every interval for direct editing.

What the editor fixes. The pass deletes stair-climb and driving-induced intervals the detector mistook for rides, splits intervals where two consecutive rides merged across a short pause, and retags any direction the velocity-sign rule flipped from a low-amplitude pre-ride drift. Dragging a mis-aligned endpoint on the plot snaps it onto the altitude step; the editor commits the GT CSV on save. Every ground-truth file in the dataset has at least one human-confirmed pass.

B.3 Buildings and phones

The campaign covers seven elevators across six buildings in three Israeli cities — Ra’anana, Herzeliya, and Ramat Gan — yielding 195 labelled rides across eight sessions. Six of the eight sessions ran the standard four-phone fleet (two experimenters $\times$ two phones each): Google Pixel 10, Samsung Galaxy A23, Samsung Galaxy S23, and Xiaomi Redmi Note 12. The Haari 3 and Bar-Ilan 2 sessions deviated to a single barometer-equipped phone (Samsung Galaxy Z Flip6 and Google Pixel 10 respectively). Ride counts per session range from 10 (Millennium Hotel exp1) to 44 (Bar-Ilan 2); Table 2 lists every session with its phones, latitude and longitude where surveyed, building type, and ride count. Beit Yitzchaki B is the held-out blind-test building; the other seven sessions feed training.

A companion data drop at [TODO: <DRIVE-URL>] mirrors the complete raw sensor logs, gramushka tables, per-session ground-truth CSVs, and editor screenshots for reproducibility.

Session	Type	Lat, Long	Phones	Rides
Haari 3, Ramat Gan, 2026-04-10	Residential	32.0658, 34.8233	Z Flip6	39
Bar-Ilan 2, Herzeliya, 2026-03-24	Residential	32.166394, 34.842190	Pixel 10	44
Millennium Hotel, exp1, 2026-04-15	Hotel	—	Pixel 10, Note 12, A23, S23	10
Millennium Hotel, exp2, 2026-04-15	Hotel	—	Pixel 10, A23, S23, Note 12	31
Millennium Outside, exp3, 2026-04-15	Hotel (external)	—	Pixel 10, A23, S23, Note 12	12
Acro Building, exp4, 2026-04-15	Office	—	Pixel 10, A23, S23, Note 12	12
Beit Mansour 1, exp5, 2026-04-15	Residential	—	Pixel 10, A23, S23, Note 12	18
Beit Yitzchaki B, exp6, 2026-04-15	Residential (test)	—	Pixel 10, A23, S23, Note 12	29

Table 2. The eight recording sessions of the field campaign. Phone short names: Pixel 10 = Google Pixel 10, A23 = Samsung Galaxy A23 (SM-A235F), S23 = Samsung Galaxy S23 (SM-S911B), Note 12 = Xiaomi Redmi Note 12 (22101320I), Z Flip6 = Samsung Galaxy Z Flip6. Latitude and longitude are recorded for two buildings; the remaining four are kept anonymised pending publication.

C Segmentation: matched-filter and pair-filter details

This appendix derives the matched-filter solution and the shared-shape joint score of §4.1, then walks every calibrated threshold of Table 3 row by row. Each subsection takes one row of the table, explains the intuition behind the test, and — where the picture is worth more than the prose — carries a single illustrative figure.

Stage / gate	Condition	Threshold
Detect: shape floor	$R^2(i) \geq$	0.40
Detect: amplitude floor	$ \hat{A}(i) \geq$	0.20 m/s ²
Detect: peak dedup	non-max suppression radius	1.0 s
Detect: same-sign gap	min. gap, equal signs	5.0 s
Gate 1: temporal gap	$t(i_2) - t(i_1) \in$	$[0, 30] \text{ s}$
Gate 2: joint score	$\frac{1}{2}(R_1^2 + R_2^2) \geq$	0.90
Gate 3: shared amplitude	$A^* \geq$	0.30 m/s ²
Gate 4: grid energy	$E(i, j) \geq$	0.40
Gate 5: quiet middle	$\text{RMS}_{\text{mid}}/A^* \leq$	0.5
Gate 6: duration penalty	$\text{rank} = S - \lambda \Delta t$	$\lambda = 0.01 \text{ s}^{-1}$

Table 3. Calibrated detection and pair-filter thresholds. The kinematics fix the template-grid bounds; a grid search on the training experiments chose the acceptance thresholds against a composite segmentation metric. The values shown are the detector’s deployed defaults.

C.1 Matched-filter solution and template grid

Before the row-by-row walk, we collect the three pieces of machinery the gates all rely on: the per-window least-squares fit, the (W, f) template grid that fit ranges over, and the shared-shape joint score that turns two candidate lobes into one pair fit.

Matched-filter least-squares solution. The detector fits each window $a_{[i]}$ with a single scaled template, minimising $\|a_{[i]} - A \tau_{W,f}\|^2$ over A . The normal equation $A \| \tau_{W,f} \|^2 = \langle a_{[i]}, \tau_{W,f} \rangle$ gives the fitted amplitude and the coefficient of determination of (13). The score $R_{W,f}^2(i)$ is the window-energy fraction the template explains: $R^2 = 1$ an exact fit, $R^2 \leq 0$ no better than the zero signal.

Template grid. The kinematics of §2.5 fix the grid. The commercial parameter ranges of §2.2 map through (10)–(11) to a half-width $W \in [0.8, 2.5] \text{ s}$ and a flat fraction $f \in [0, 0.7]$; the grid pads these to 30 log-spaced W in $[0.3, 3.0] \text{ s}$ and 15 linear f in $[0.05, 0.80]$, with an $f=0$ triangle row for short rides that never reach the a_{max} plateau.

Shared-shape joint score. The pair score (17) fits one amplitude to both lobes of a candidate pair. With a_k the k th lobe window, $P_k = \|a_k\|^2$, $u_k = s_k \langle a_k, \tau_{W,f} \rangle$ the sign-corrected correlation, and $N = \|\tau_{W,f}\|^2$, minimising $\sum_k \|a_k - s_k A \tau_{W,f}\|^2$ over the shared A gives $A^* = (u_1 + u_2)/(2N)$ and the per-lobe R_k^2 of (17). Unrelated bursts peak at a different cell for each lobe, so forcing one (W, f) on both is what makes the score discriminating.

C.2 Detection shape and amplitude floors

The first two rows of Table 3 reject samples that cannot be a take-off or landing lobe at all (15). The shape floor $r_{\text{min}} = 0.40$ is deliberately permissive: hand-held jitter and clipped windows pull a real lobe’s R^2 below a textbook trapezoid’s, so

the detector leaves the strict shape test to the pair stage. The amplitude floor $a_{\min} = 0.20 \text{ m/s}^2$ drops lobes buried in sensor and walking noise, and rises per phone when a device's measured noise runs higher. Both favour recall — a missed lobe is lost for good, while a spurious one still faces the pairing gates.

C.3 Peak de-duplication and same-sign gap

The next two rows of Table 3 stop one physical lobe from entering the pair stage as several peaks. A lobe's matched-filter score forms a ridge several samples wide, so non-maximum suppression over a 1.0 s radius keeps only its crest, and the same-sign gap collapses any two *equal*-sign peaks closer than 5.0 s — a lobe re-detected at a neighbouring W . Opposite-sign peaks pass untouched: separating a take-off from its landing is the pairing stage's task.

C.4 Gate 1 — pairing window $[t_{\min}, t_{\max}]$

Gate 1 bounds how far a take-off and a landing may sit apart before the pair is scored (16) — a search window, not a model of ride duration. The lower bound $t_{\min} = 0 \text{ s}$ admits the shortest single-floor hops, whose lobes almost touch; the upper bound $t_{\max} = 30 \text{ s}$ brackets one uninterrupted ride at the car speeds and inter-floor heights of vertical-transportation practice [4, 12]. A high-rise express run would need a wider window; the 30 s cap trades that reach against joining lobes drawn from *different* rides.

C.5 Gate 2 — shared-shape joint score r_{pair}

Gate 2 demands a joint score of at least $r_{\text{pair}} = 0.90$, far above the 0.40 detector floor. A genuine ride offers a constraint a single lobe never does: both lobes are cut from *one* cabin profile, so one (W, f) template must fit them together (17). Unrelated bursts each fit some template alone but peak at different cells, and no shared cell scores both above 0.90. The strict floor is cheap — it costs almost no true rides while removing the bulk of accidental pairings.

C.6 Gate 3 — shared amplitude a_{pair}

Gate 3 requires the shared amplitude A^* to clear $a_{\text{pair}} = 0.30 \text{ m/s}^2$, above the 0.20 detector floor. Since A^* is the mean of the two lobe amplitudes (17), a pair survives only when *both* lobes are firm: a strong take-off matched to a weak noise lobe averages below 0.30. The gate thus rejects the asymmetric pairings the per-lobe floor lets through.

C.7 Gate 4 — grid energy $e_{\min}$

Gate 4 measures *how broadly* the grid supports the match. A genuine ride near-fits a whole neighbourhood of (W, f) templates — a wide bright band in the joint- R^2 heatmap — while a coincidental pairing fits one lucky cell and leaves the grid dark. Grid energy E averages the clamped joint R^2 over every cell (19); the floor $e_{\min} = 0.40$ sits above a single-cell match and below any broad band, catching false pairs that clear the best-cell score of Gate 2 (Figure 14).




Fig. 14. The grid-energy gate $e_{\min}$. Each panel is the joint- R^2 heatmap over the (W, f) template grid for one candidate pair. A genuine ride (left) is a near-fit across a broad band of templates, so its grid energy E sits well above the 0.40 floor; a coincidental pairing (right) fits one narrow sliver and leaves the grid dark, so its energy falls far below the floor and Gate 4 rejects it.

C.8 Gate 5 — quiet-middle cruise gate ρ

Between its two lobes a real ride holds constant velocity, so the vertical acceleration there is near zero. Gate 5 takes RMS_{mid} , the root-mean-square of the smoothed acceleration on the inter-lobe window $(t_{c_1} + W, t_{c_2} - W)$, and rejects the pair when it exceeds ρA^* with $\rho = 0.5$ — when the cruise wobble passes half the lobe amplitude. The test catches pairs whose middle never settles, most often a walking false positive where continuous motion keeps the inter-lobe stretch as loud as the lobes (Figure 15).




Fig. 15. The quiet-middle gate ρ . **Left:** a genuine ride — the smoothed acceleration returns to zero on the inter-lobe cruise (shaded), so $\text{RMS}_{\text{mid}}/A^*$ stays well under $\rho = 0.5$ and the pair is accepted. **Right:** a pair whose inter-lobe stretch never goes quiet — $\text{RMS}_{\text{mid}}/A^*$ exceeds 0.5 and Gate 5 rejects it.

C.9 Gate 6 — duration penalty λ and greedy resolution

A passenger often rides in stages — floor 1 to 4, a stop, then 4 to 8 — recorded as two back-to-back rides and *four lobes* (Figure 16). The take-off of the first ride and the landing of the last form a *super pair*: it shares the same cabin shape, so its joint score S matches either true ride and Gate 2 cannot tell them apart. Gate 6 breaks the tie by ranking every admissible pair by $S - \lambda \Delta t$ with $\lambda = 0.01 \text{ s}^{-1}$ (20): the penalty docks the long span far more than the short ones, so the two short rides out-rank the super pair. A greedy resolver commits pairs highest-first, skipping any that reuse a committed lobe, and the super pair never commits.




Fig. 16. The duration penalty λ on an up-going multi-stop trip. **Top:** the acceleration trace — four lobes, two back-to-back rides (shaded). **Bottom:** the matched-filter signed- R^2 score, peaking at all four lobes. The take-off of the first ride and the landing of the last form a *super pair* whose joint score S matches either short ride; only the rank $S - \lambda \Delta t$ separates them, docking the long span hardest so the two short rides commit first.

D Prediction: error budget and conformal calibration

This appendix derives the closed-form height (23), the per-algorithm error budgets of §4.2.3, and the conformal coverage guarantee of (28), and it explains and illustrates every quality-filter check of §4.2.5.

D.1 The trapezoid pulse-pair integral

The closed-form height (23) follows from one elementary integral, the area of the unit trapezoid. The trapezoid $\tau_{W,f}$ of (9) is a plateau of width $2fW$ joined to zero by two ramps of base $(1-f)W$, so it integrates to

$$\int_{-\infty}^{\infty} \tau_{W,f}(t) dt = \underbrace{2fW}_{\text{plateau}} + 2 \cdot \underbrace{\frac{1}{2}(1-f)W}_{\text{two ramps}} = W(1+f). \quad (29)$$

A single acceleration lobe $sA\tau_{W,f}$ integrates once to a velocity that ramps the cabin from rest to the cruise speed $v_{\text{peak}} = AW(1+f)$, the lobe area of (29). The take-off lobe ramps the cabin up to v_{peak} ; the opposite-sign landing lobe brings it back to rest; between them the cabin holds v_{peak} . Because $\tau_{W,f}$ is symmetric about its centre, the velocity rises point-symmetrically about each lobe centroid, so each $2W$ -wide ramp window contributes the same area as a width- W stretch at full v_{peak} . Integrating the velocity over the ride therefore gives

$$\Delta h = \int_0^T v(t) dt = s v_{\text{peak}} \Delta t_c = s A W (1+f) \Delta t_c,$$

which is (23): the cruise velocity times the centroid spacing, signed by the ride direction.

Joined-pulse short-ride variant. A short hop has no cruise: the two lobes touch at $\Delta t_c = 2W$. Setting $\Delta t_c = 2W$ in (23) gives

$$\Delta h = 2sAW^2(1+f), \quad (30)$$

the joined-pulse height. The arrival time drops out, so Δh depends on three parameters, with gradients $\partial\Delta h/\partial A = 2sW^2(1+f)$, $\partial\Delta h/\partial W = 4sAW(1+f)$, and $\partial\Delta h/\partial f = 2sAW^2$. The joined template packs twice the energy of a single lobe, $\langle\tau_{\text{joined}}, \tau_{\text{joined}}\rangle = 2\langle\tau, \tau\rangle$, so its amplitude bound is correspondingly tighter. Prediction runs the pair fit and the joined fit on every segment and keeps the higher joint R^2 , the joined fit's one fewer free parameter breaking near-ties in its favour.

Agreement with the S-curve. The trapezoid is not an approximation of the S-curve of §2.3 — it is its acceleration lobe. Under the bang-bang jerk command (3) the jerk is piecewise constant, so each acceleration lobe is piecewise linear: exactly the trapezoid $\tau_{W,f}$, with (W, f) fixed by the machine parameters through (10)–(11). A controller that caps the jerk derivative rounds the four corners of that lobe into short curved transitions. The trapezoid is then the area- and centroid-matched linearisation of the rounded lobe — the unique trapezoid with the same integral and the same first moment — so the lobe area

$$\int a_{\text{lobe}}(t) dt = AW(1+f) \quad (31)$$

and the centroid spacing Δt_c , hence the height (23), are preserved under jerk rounding (Figure 17). The fit recovers Δh without resolving the corner shape it cannot constrain.




Fig. 17. One acceleration lobe of the S-curve. Under the bang-bang jerk command (3) the lobe is exactly the trapezoid $\tau_{W,f}$ (solid), with (W, f) fixed by the machine parameters through (10)–(11). A controller that caps the jerk derivative rounds the four corners (dashed); the trapezoid is the area- and centroid-matched linearisation of that rounded lobe, so the shaded area $AW(1+f)$ — the cruise velocity — and the centroid are preserved, and with them the closed-form height (23).

D.2 ZUPT error budget

The ZUPT standard deviation (25) sums four independent variances.

White-noise double integration. Integrating a white-noise acceleration sequence of per-sample standard deviation σ_a once produces a velocity random walk; integrating again produces a position. After the linear drift term of (21) removes the unconstrained ramp mode, the end-point position variance is $\sigma_a^2 \Delta t^4 N^3 / 12$ over N active samples, whose root is (24). The cube in N is the double-integration random-walk scaling; the factor $1/12$ is the residual after the drift correction, against the $1/3$ of an uncorrected double integral.

Mechanical jitter. A hand-held phone is not a rigid inertial platform: hand tremor, pocket friction, and skin-coupled breathing add a coloured-noise acceleration band the datasheet omits. The budget models it as an extra acceleration $\sigma_{\text{mech}}^a = j_0 (1 + \min(\rho/10^\circ, 2))$, a baseline $j_0 = 0.03 \text{ m/s}^2$ scaled up to threefold by the ride drift angle ρ , then integrated under ZUPT exactly as the white-noise term: $\sigma_{\text{mech}} = \sigma_{\text{mech}}^a \Delta t^2 \sqrt{N^3/12}$.

Relative scale. The longer the ride, the larger the absolute error, because the unmodelled coloured drift is roughly linear in the path integral of acceleration. A relative floor $\sigma_{\text{rel}} = k_{\text{rel}} |\Delta h|$ with $k_{\text{rel}} = 0.05$ captures it.

Projection penalty. A ride that lacks a clean stationary window projects its vertical axis from a one-sided or in-ride gravity estimate, losing one or two degrees of freedom on the vertical direction. The penalty $\sigma_{\text{proj}} = (\pi - 1) \cdot 0.25 |\Delta h|$ scales with a method factor $\pi \in \{1.0, 1.2, 1.3, 3.0\}$ for the pre-and-post, pre-only, post-only, and magnitude-only projections, further multiplied by $1 + \min(\rho_{\text{pp}}/12.5^\circ, 2)$ for the pre/post rotation angle ρ_{pp} . The four terms combine in quadrature (25), floored at $\sigma_{\text{ZUPT}} = \max(\sqrt{\Sigma}, 0.15 \text{ m})$ so the conformal multiplier cannot blow up on a degenerate fit.

D.3 Trapezoid error budget

Matched-filter Cramér–Rao bound. Model the vertical trace as the fitted template pair plus Gaussian noise of standard deviation σ_a^{eff} . The Fisher information of this likelihood bounds the variance of each fitted parameter. For the amplitude and the lobe centre,

$$\sigma_A^2 = \frac{(\sigma_a^{\text{eff}})^2}{\langle \tau, \tau \rangle}, \quad \sigma_{t_c}^2 = \frac{(\sigma_a^{\text{eff}})^2}{A^2 \langle \dot{\tau}, \dot{\tau} \rangle}, \quad \sigma_{\Delta t_c}^2 = 2 \sigma_{t_c}^2, \quad (32)$$

where $\langle \tau, \tau \rangle$ is the template energy of (13) and $\langle \dot{\tau}, \dot{\tau} \rangle = 2/[(1-f)W]$ is the slope energy of the two ramps. A wider, sharper-edged lobe carries more energy and more slope, so it pins both amplitude and timing more tightly. The grid quantises the shape parameters, so their variance is that of a uniform prior over one grid cell, $\sigma_W^2 = \Delta_W^2/12$ and $\sigma_f^2 = \Delta_f^2/12$.

Delta-method propagation. Differentiating the closed-form height (23) gives the four gradients

$$\frac{\partial \Delta h}{\partial A} = sW(1+f)\Delta t_c, \quad \frac{\partial \Delta h}{\partial W} = sA(1+f)\Delta t_c, \quad \frac{\partial \Delta h}{\partial f} = sAW\Delta t_c, \quad \frac{\partial \Delta h}{\partial \Delta t_c} = sAW(1+f), \quad (33)$$

and substituting (32) and (33) into (26) yields the per-segment variance $\sigma_{\Delta h}^2$.

Data-adaptive effective noise. The white-noise figure from the sensor datasheet understates the error on a hand-held phone by an order of magnitude. The pipeline replaces it with the residual root-mean-square of the fitted template,

$$\sigma_a^{\text{eff}} = \max\left(\sqrt{\frac{1}{N} \|a - \hat{a}_{\theta^*}\|^2}, \sigma_a^{\text{ds}}\right),$$

floored at the datasheet value σ_a^{ds} . Whatever the template fails to explain — tremor, projection wobble, coloured drift — enters the residual, so the budget counts it automatically. A second correction divides the effective variance by the fit quality, $\tilde{\sigma}_a^2 = (\sigma_a^{\text{eff}})^2 / \max(R_{\text{joint}}^2, \varepsilon)$ with $\varepsilon = 10^{-3}$ — the Wald post-regression form, in which a poorly fitting template widens the posterior on every parameter in proportion.

Velocity-anchored amplitude variance. When the fit of §4.2.2 replaces the matched-filter amplitude with the measured cruise velocity, the CRB (32) no longer describes σ_A . The cruise velocity is a mean of white-noise velocity samples, so $\sigma_v^2 = \tilde{\sigma}_a^2 \Delta t / T_{\text{cruise}}$ and $\sigma_A^2 = \sigma_v^2 / [W(1 + f)]^2$, which the pipeline substitutes for σ_A^2 whenever the cruise window is long enough to anchor ($T_{\text{cruise}} > 0.5$ s).

Overlap inflation of the centroid spacing. The CRB (32) treats the two lobe centres as independent. At $\Delta t_c = 2W$ the lobes touch but their supports stay disjoint, the off-diagonal Fisher block is zero, and the diagonal CRB is exact. For $\Delta t_c < 2W$ — unphysical for the shared-shape pair model — the two centre estimates correlate, and the budget inflates the spacing variance by

$$\sigma_{\Delta t_c}^2 = 2 \sigma_{t_c}^2 \left[1 + \left(\frac{2W - \Delta t_c}{2W \delta} \right)^2 \right], \quad \delta = 0.05, \quad (34)$$

and the factor is 1 in the physical regime $\Delta t_c \geq 2W$. The inflation keeps the budget well-conditioned without penalising a genuine touching-lobe ride, which the joined-pulse fit (30) handles directly.

D.4 Split-conformal coverage and calibration

The multiplier k of §4.2.4 is the $\lceil (1 - \alpha)(n + 1) \rceil / n$ empirical quantile of the n calibration nonconformity scores (27). For exchangeable rides the rank of a test score among the calibration scores is uniform, so the test score falls at or below that quantile with probability at least $1 - \alpha$ — which is (28), with no assumption on the shape of the error distribution.

The two algorithms calibrate k on different pools. ZUPT calibrates on every clean ride: the reported metrics are computed on clean data regardless of the quality verdict, so a consumer who overrides the filter still receives the promised coverage. The trapezoid fit calibrates on the rides that are both clean and quality-accepted — its exchangeable population, since a rejected trapezoid fit carries a σ the budget knowingly under-predicts. Exchangeability fails across the clean/corrupted split, so neither pool draws on corrupted recordings. The pipeline floors the multiplier at the Gaussian 90 % half-width 1.645, and keeps an absolute floor and cap on the reported $k\sigma$ as a guard against a pathologically small or large fitted σ .

D.5 Quality-filter checks

Each algorithm runs its own quality filter (§4.2.5). Both share the same machinery: hard gates that reject outright at the first breach, and graded score terms that sum into a penalty rejecting the ride above 6. Table 4 lists the ZUPT checks and Table 5 the trapezoid checks. The remaining paragraphs give the intuition behind each feature and a real-recording good-versus-bad example.

Gravity and orientation. Any accelerometer estimate is only as good as its vertical axis, so both filters open with three orientation checks. Pre-window stationarity s_{pre} is the spread of the gravity vector over the bracketing rest window: a still phone holds three flat axis traces, and a window the filter cannot certify forces the magnitude fallback and a rejection (Figure 18). The pre/post tilt is the angle between the gravity vectors of the two rest windows; past 25° the phone was reoriented mid-ride and the single projection no longer holds for both ends (Figure 19). Ride drift is the

Check	Condition	Threshold
<i>Hard gates</i>		
Gravity calibration	signed vertical projection available	reject if magnitude-only
Phone tilt	$\angle(g_{\text{pre}}, g_{\text{post}}) \leq$	25°
Ride drift	max gravity drift $\leq$	15°
Impact peak	$\max a - \bar{a} \leq$	8 m/s^2
Active motion	$\text{fraction } a > 0.10 \text{ m/s}^2 \geq$	0.04
Displacement	$ \Delta h \in$	$[0.4, 100] \text{ m}$
Segment length	active samples $N \geq$	30
<i>Score terms</i> (penalty sum rejects above 6)		
Pre-window stability	s_{pre}	graded
Pre/post tilt	$\angle(g_{\text{pre}}, g_{\text{post}})$	graded from 10°
Ride drift	max gravity drift	graded
Impact peak	$\max a - \bar{a} $	graded at 4, 6 m/s^2
End-velocity ratio	$ v(T) /\max v $	graded

Table 4. ZUPT quality-filter checks and calibrated thresholds.

Check	Condition	Threshold
<i>Hard gates</i>		
Gravity calibration	signed vertical projection available	reject if magnitude-only
Phone tilt	$\angle(g_{\text{pre}}, g_{\text{post}}) \leq$	25°
Joint R^2 floor	$R_{\text{joint}}^2 \geq$ (short/mid/long)	0.35/0.50/0.60
Displacement	$ \Delta h \in$	$[0.4, 120] \text{ m}$
Lobe spacing	$\Delta t_c \geq$	0.2 s
Active motion	$(2W + \Delta t_c)/T_{\text{ride}} \geq$	0.05
Amplitude anchor	$A_{\text{used}}/A_{\text{fit}} \in$	$[0.5, 2.0]$
Out-of-lobe residual	residual density, outside/inside $\leq$	4
Cruise steadiness	$\text{CV}(v_{\text{cruise}}) \leq$	0.6
Segment length	samples $N \geq$	30
<i>Score terms</i> (penalty sum rejects above 6)		
Pre-window stability	s_{pre}	graded
Pre/post tilt	$\angle(g_{\text{pre}}, g_{\text{post}})$	graded from 10°
Ride drift	max gravity drift	graded
Joint R^2	R_{joint}^2	graded below 0.8
Residual structure	lag-1 residual ACF	graded at 0.25, 0.4, 0.6

Table 5. Trapezoid quality-filter checks and calibrated thresholds. The filter accepts a ride only when every hard gate passes and the penalty score stays below its cap.

largest angle the gravity direction sweeps across one-second chunks of the ride itself; a held phone stays within a few degrees, while active handling climbs past the 15° gate (Figure 20).

Impact (ZUPT). ZUPT adds an impact gate: the largest deviation of the acceleration magnitude from its mean. A tap, a pocket drop, or a phone hand-off spikes the trace far past the elevator’s own sub-metre-per-second-squared lobes, and double integration would fold that spike straight into Δh . The gate rejects past 8 m/s^2 and the score term grades the approach to it (Figure 21). Two further ZUPT checks keep only soft score terms. No labelled ride trips their hard gates: every recorded ride carries ample motion, so the active-motion gate never fires, and the magnitude-detrended trace drives the end-velocity ratio close to zero by construction.




Fig. 18. Pre-window stationarity s_{pre} . **Left:** a still bracketing window — three flat axis traces — certifies the gravity calibration. **Right:** a window the filter cannot certify; the projection falls back to the unsigned magnitude and the ride is rejected.




Fig. 19. Pre/post gravity-vector angle. **Left:** the gravity components hold level across the two rest windows, so one projection serves both ride ends. **Right:** the components jump between windows — the phone was reoriented mid-ride — and the tilt clears the 25° gate.

Fit quality (trapezoid). The trapezoid filter adds two checks that the fitted template actually explains the trace. The joint R^2 is the matched-filter fit quality across both lobes; its floor is graded by predicted displacement — 0.35/0.50/0.60 for short, mid, and long rides — because a short ride's signal-to-noise ratio is intrinsically worse (Figure 22). The residual autocorrelation scores what the R^2 misses: a template that captures the lobe energy but leaves a slow, structured residual signals an unmodelled disturbance. On the smoothed trace adjacent residuals are correlated by construction, so the score knees sit high, at 0.25, 0.4, and 0.6 (Figure 23).

Fitted geometry (trapezoid). Three checks bound the fitted geometry against matched-filter misfires. The amplitude anchor ratio A_{used}/A_{fit} compares the velocity-anchored amplitude with the raw matched-filter amplitude; a trustworthy fit needs little rescaling, and a ratio outside [0.5, 2.0] flags a template that locked onto the wrong scale (Figure 24). The




Fig. 20. In-ride gravity drift. **Left:** the gravity direction holds within a few degrees through the ride. **Right:** the direction sweeps past the 15° gate, so no single vertical axis describes the whole segment.




Fig. 21. Impact-peak magnitude. **Left:** the magnitude-detrended trace stays within the elevator's own acceleration band. **Right:** a hand-off spike clears the 8 m/s² gate; double integration would fold it into the height.

out-of-lobe residual ratio compares the residual energy density outside the two lobe windows with the density inside; a healthy fit keeps the outside density no higher than the inside, and a ratio past fourfold signals a mid-ride disturbance the matched filter could not explain (Figure 25). The cruise coefficient of variation is the spread of the ZUPT-integrated velocity across the inter-lobe window; a genuine cruise is a flat velocity plateau, and a coefficient past 0.6 means the cabin never reached steady cruise (Figure 26).




Fig. 22. Matched-filter joint R^2 . **Left**: the fitted trapezoid pair hugs the smoothed trace. **Right**: a trace with no clean pulse structure leaves the template far short of the graded R^2 floor.




Fig. 23. Lag-1 residual autocorrelation. **Left**: the fit residual scatters tightly around zero. **Right**: the residual carries a slow, coherent excursion — structure the trapezoid template never modelled.




Fig. 24. Velocity-anchor amplitude ratio. **Left:** the velocity-anchored amplitude matches the matched-filter amplitude, so the two templates coincide. **Right:** the anchor rescales the amplitude far past the $[0.5, 2.0]$ band — a matched-filter misfit.




Fig. 25. Out-of-lobe residual concentration. **Left:** the residual energy stays within the two lobe windows. **Right:** residual energy concentrates outside the two lobe windows — a mid-ride disturbance the template never modelled.




Fig. 26. Cruise-velocity coefficient of variation. **Left:** the ZUPT-integrated velocity holds a flat plateau between the lobes. **Right:** the plateau wobbles past the 0.6 gate — the cabin never settled into steady cruise.